



#FOSDEM #MySQLDay #MySQL8isGreat



Friends let real friends use MySQL 8.0



30th January 2020
BRUSSELS

Saverio Miroddi
Senior Engineer at  ticketsolve



{ FRIENDS DON'T LET
REAL FRIENDS USE
RELATIONAL DATABASES }



mongoDB.

14' x 48'

004000

CLEAR CHANNEL

#15277

FRIENDS ~~DON'T~~ LET
REAL FRIENDS USE
~~RELATIONAL DATABASES~~

MYSQL



14' x 48'

CLEAR CHANNEL

#15277

Introduction: Who are I and Ticketsolve

- ▶ Lead of invisible things at Ticketsolve
- ▶ Ticketsolve is a primary player in Ireland/UK in the events ticketing market, and it has an internal open source program
- ▶ Efficiency is one of our key values; this reflect in the presentation's tooling and general approach

Introduction: Presentation and target audience

- ▶ We've been using MySQL for 10+ years
- ▶ We starting each migration process around 6 months after GA
- ▶ This presentation describes our experience in migrating to MySQL 8.0; it covers the most important concepts of what a sysadmin needs to know and do, if "tomorrow they need to migrate production to 8.0"
- ▶ The most important subject is utf8mb4 (and its default collation)
- ▶ There is no target level, although the audience is system/database administrators

Preparing MySQL: setup and tooling

- ▶ Use tarballs for convenience, and a fixed path with symlink
- ▶ I'm using several scripts to speed up operations and access frequently used terms
- ▶ PS. I'm running fully on ZFS!!

First step: compare the global system variables

- ▶ Although we read the release notes for every release (including patch versions), approaching major version changes is very intimidating
- ▶ Comparing system variables is a good starting point
- ▶ Check out the bulky system variables, and filter them out

Migrating to utf8mb4: Summary

- ▶ The new collation is the crucial aspect to consider
- ▶ What ``utf8mb4_0900_ai_ci`` is and what it brings
- ▶ Unfortunately, the documentation available in the web is outdated - it used the old default collation (``utf8mb4_general_ci``)

Example of real-world issues:

- ▶ After migrating to ``utf8mb4_general_ci``, it's expensive to move to ``utf8mb4_0900_ai_ci``
- ▶ Untrimmed strings caused unexpected results in queries

How the charset parameters work

- ▶ What happens when a client connects to the server?
- ▶ Which are the database object charset-related defaults?

Collation coercion, and issues ``general` <> `0900_ai``

I'll simulate all the cases of inter-charset comparisons.

- ▶ Comparisons ``utf8_general_ci` columns <> literals` (see next)
- ▶ Comparisons ``utf8_general_ci` columns <> columns`
 - ▶ Moving a schema from ``utf8mb4_general_ci`` to ``utf8mb4_0900_ai_ci`` is not trivial!

Collation Coercibility in Expressions

- 0: An explicit COLLATE clause (not coercible at all)
- 1: The concatenation of two strings with different collations
- 2: The collation of a column or a stored routine parameter or local variable
- 3: A “system constant” (the string returned by functions such as USER() or VERSION())
- 4: The collation of a literal
- 5: The collation of a numeric or temporal value
- 6: NULL or an expression that is derived from NULL

(From <https://dev.mysql.com/doc/refman/8.0/en/charset-collation-coercibility.html>)

Issues with `0900_ai` collation padding

- ▶ How are trailing whitespaces handled?
- ▶ What defines such handling? (see next)

The following are the core string comparison rules from the SQL (2003) standard (section 8.2):

3) The comparison of two character strings is determined as follows:

a) Let CS be the collation [...]

b) **If the length in characters of X is not equal to the length in characters of Y, then the shorter string is effectively replaced, for the purposes of comparison, with a copy of itself that has been extended to the length of the longer string by concatenation on the right of one or more pad characters**, where the pad character is chosen based on CS. **If CS has the NO PAD characteristic, then the pad character is an implementation-dependent character** different from any character in the character set of X and Y that collates less than any string under CS. Otherwise, the pad character is a space.

Triggers

- ▶ Triggers are fairly easy to handle (see next), as they can be dropped/rebuilt - just make sure to consider comparisons in the trigger body

```

SHOW CREATE TRIGGER enqueue_comments_update_instance_event\G

-- SQL Original Statement:
CREATE TRIGGER `enqueue_comments_update_instance_event`
AFTER UPDATE ON `comments`
FOR EACH ROW
trigger_body: BEGIN
    SET @changed_fields := NULL;

    IF NOT (OLD.description <=> NEW.description COLLATE utf8_bin AND CHAR_LENGTH(OLD.description) <=> CHAR_LENGTH(NEW.description)) THEN
        SET @changed_fields := CONCAT_WS(',', @changed_fields, 'description');
    END IF;

    IF @changed_fields IS NOT NULL THEN
        SET @old_values := NULL;
        SET @new_values := NULL;

        INSERT INTO instance_events(created_at, instance_type, instance_id, operation, changed_fields, old_values, new_values)
        VALUES(NOW(), 'Comment', NEW.id, 'UPDATE', @changed_fields, @old_values, @new_values);
    END IF;
END

-- character_set_client: utf8mb4
-- collation_connection: utf8mb4_0900_ai_ci
-- Database Collation: utf8mb4_0900_ai_ci

```


Behavior with indexes

- ▶ How does an ``utf8mb4`` index behave when comparing with an ``utf8`` literal?
- ▶ How does it behave on joins? What does one see in the standard and extended ``EXPLAIN``s?

Consequences of the increase in (potential) size of char columns

- ▶ utf8mb4 characters will take 33% more, which must stay within the InnoDB index limit, which is however (as of 8.0 default), high (3072 bytes).

Remember:

- ▶ `[VAR]CHAR(n)` refers to the number of characters; therefore, the maximum requirement is $4 * n$ bytes
- ▶ `TEXT` fields refer to the number of bytes

Mac Homebrew default collation is ``utf8mb4_general_ci``!

- ▶ When MySQL is installed via Homebrew, the default collation is ``utf8mb4_general_ci``.
- ▶ I've opened an issue and sent a PR
- ▶ Interim solution 1: edit the formula and rebuild mysql
- ▶ Interim solution 2: disable the server handshake

`information_schema_stats_expiry` introduction

- ▶ When querying the `information_schema`, beware that stats are updated only once a day!
- ▶ If required, invoke `ANALYZE TABLE` (or reduce the setting value)

GROUP BY is now unsorted (not implicitly sorted)

- ▶ Developers will need to go through the codebase and check each query with GROUP BY
- ▶ It's impossible to use a one-size-fits-all approach

Very generic approaches:

- ▶ `grep -A`
- ▶ mini script checking every pair of lines

Schema migration tool issues

- ▶ There's a known [showstopper bug](<https://github.com/github/gh-ost/issues/687>) on the latest Gh-ost release, which prevents operations from succeeding on MySQL 8
 - ▶ There's a test fix PR, it would be great if somebody with scale could test it!
- ▶ Use trigger-based tools, like ``pt-online-schema-change`` v3.1.1 or v3.0.x (but v3.1.0 is broken!) or Facebook's OnlineSchemaChange

Conclusion

- ▶ Over the next weeks, I will expand this subject into a series of articles in my professional blog
- ▶ This presentation is hosted (with git history preserved) at <https://github.com/saveriomiroddi/prefosdem-2020-presentation>
- ▶ Those curious about ZFS can check out my installer: <https://github.com/saveriomiroddi/zfs-installer>
- ▶ Questions!